

An EKF-SLAM toolbox in Matlab

Joan Solà
LAAS-CNRS

March 26, 2009

Contents

1	The SLAM toolbox presentation	1
2	Data organization	4
2.1	SLAM data	4
2.2	Simulation data	9
2.3	Graphics data	13
2.4	Plain data	16
3	Functions	17
3.1	High level	17
3.2	Interface level	17
3.3	Low level library	20
4	Developing new models	20

1 The SLAM toolbox presentation

In a typical SLAM problem, one or more robots navigate an environment, discovering and mapping landmarks on the way by means of their onboard sensors. Observe in Fig. 1 the existence of robots of different kinds, carrying a different number of sensors of different kinds, and observing landmarks of different kinds. All this variety of data is handled by the present toolbox in a way that is quite transparent.

In this toolbox, we organized the data into three main groups, see Table 1. The first group contains the objects of the SLAM problem itself, as they appear in Fig. 1. A second group contains objects for simulation. A third group is designated for graphics output.




Figure 1: Overview of the SLAM problem with the principal data structures.

Table 1: All data structures.

Purpose	SLAM	Simulator	Graphics
Robots	Rob	SimRob	
Sensors	Sen	SimSen	SenFig
Landmarks	Lmk	SimLmk	
Observations	Obs	SimObs	
Map	Map		MapFig
Time	Tim		

Apart from the data, we have of course the functions. We organized them in three levels, from most abstract and generic to the basic manipulations, as is sketched in Fig. 2. The highest level deals exclusively with the structured data we mentioned just above, and calls functions of an intermediate level called the interface level. The interface level functions split the data structures into more mathematically meaningful elements, check data types to decide on the applicable methods, and call the basic functions that constitute the basic level.

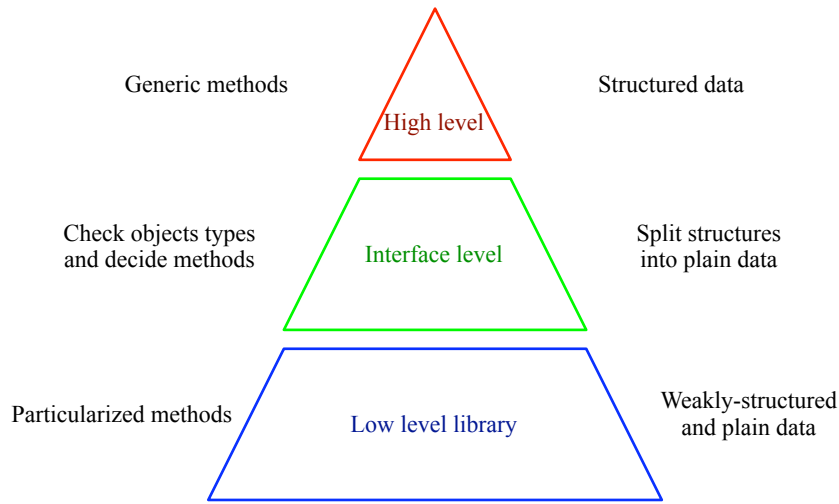


Figure 2: Overview of the levels of abstraction of the functions and their relation to data structuration. Functions and scripts in the High and Interface levels are in the `%SLAMtoolbox/HighLevel/` and `%SLAMtoolbox/InterfaceLevel/` directories. The Low Level library occupies all other directories.

2 Data organization

It follows a brief explanation of the SLAM data structures, the Simulation and Graphic structures, and the plain data types.

2.1 SLAM data

For a SLAM system to be complete, we need to consider the following parts:

Rob: A set of robots.

Sen: A set of sensors.

Lmk: A set of landmarks.

Map: A stochastic map containing the states of robots, landmarks, and eventually sensors.

Obs: The set of landmark observations made by the sensors.

Tim: A few time-related variables.

Our toolbox considers these objects as the only existing data. In Matlab, these objects are defined as structures holding a variety of fields. For the SLAM objects we used 3-letter names as in the list just above, {**Map**, **Rob**, **Sen**, **Lmk**, **Obs**, **Tim**}. The structures are reproduced in Figs. 3 to 7 for reference. Structure arrays hold any number of such objects. For example, all the data related to robot number 2 is stored in **Rob(2)**. To access the rotation matrix defining the orientation of this robot we simply use

Rob(2).frame.R

A special case I want to mention here is **Obs**, because observations relate sensors to landmarks: the data associated to the observation of landmark **lmk** from sensor **sen** is stored in **Obs(sen,lmk)**. Before reading on, please revisit Fig. 1 to see how simple things are.

It follows a reproduction of the arborescences of the principal structures in the SLAM data.

```

Rob(r)          Robot structure, containing:
  .rob          * index in Rob() array
  .id           * robot id
  .name         * robot name
  .type         * robot type
  .sensors      * list of installed sensors
  .motion       * motion model
  .con          * control structure
    .u          - control signals for the motion model
    .U          - covariance of u
  .frame        * frame structure, containing:
    .x          - 7-vector, position and orientation  $x = [t;q]$ 
    .P          - covariances matrix of x
    .t          - position
    .q          - orientation quaternion
    .R          - rotation matrix,  $R = q2R(q)$ 
    .Rt         - transposed R
    .it         - inverse position,  $it = -Rt*t$ 
    .iq         - inverse quaternion,  $iq = q2qc(q)$ 
    .Pi         - PI matrix,  $Pi = q2Pi(q)$ 
    .Pc         - conjugate PI matrix,  $Pc = pi2pc(Pi)$ 
    .r          - range in the SLAM map Map
  .vel          * velocity stucture, containing
    .x          - 6-vector, linear and angular velocities
    .P          - covariances matrix of x
    .r          - range in the SLAM map Map
  .state        * state structure, containing
    .x          - robot's state vector,  $x = [frame.x;vel.x]$ 
    .P          - covariances matrix of x
    .size       - size of x
    .r          - range in the SLAM map Map

```

Figure 3: The Rob structure array with its arborescence.

```

Sen(s)          Sensor structure, containing:
  .sen          * index in Sen() array
  .id           * sensor id
  .name         * sensor name
  .type         * sensor type
  .robot        * robot it is installed to
  .frame        * frame structure, containing:
    .x          - 7-vector, position and orientation  $x = [t;q]$ 
    .P          - covariances matrix of  $x$ 
    .t          - position
    .q          - orientation quaternion
    .R          - rotation matrix,  $R = q2R(q)$ 
    .Rt         - transposed  $R$ 
    .it         - inverse position,  $it = -Rt*t$ 
    .iq         - inverse quaternion,  $iq = q2qc(q)$ 
    .Pi         - PI matrix,  $Pi = q2Pi(q)$ 
    .Pc         - conjugate PI matrix,  $Pc = pi2pc(Pi)$ 
    .r          - range in the SLAM map Map
  .par          * sensor parameters
    .k          - intrinsic params
    .d          - distortion vector
    .c          - correction vector
    .imSize     - image size
  .state        * state structure, containing
    .x          - sensor's state vector,  $x = frame.x$  or  $x = []$ 
    .P          - covariances matrix of  $x$ 
    .size       - size of  $x$ 
    .r          - range in the SLAM map Map

```

Figure 4: The `Sen` structure array with its arborescence.

Lmk(1)	Landmark structure, containing:
.lmk	* index in Lmk() array
.id	* landmark id
.type	* sensor type
.used	* landmark is used in the map
.state	* state structure, containing
.x	- landmark's state vector
.P	- covariances matrix of x
.size	- size of x
.r	- range in the SLAM map Map
.nom	* prior of non-measurable degrees of freedom
.n	- mean
.N	- covariance
.par	* other lmk parameters

Figure 5: The Lmk structure array with its arborescence.

Map	Map structure, containing:
.used	* vector of flags indicating non-free positions
.x	* state vector's mean
.P	* covariances matrix
.size	* size of the map, in number of states

Figure 6: The Map structure with its arborescence.

Obs(s,l)	Observation structure, containing:
.sen	* index to sensor in Sen() array
.lmk	* index to landmark in Lmk() array
.sid	* sensor id
.lid	* landmark id
.meas	* measurement
.y	- mean
.R	- covariance
.nom	* non-measurable degrees of freedom
.n	- mean
.N	- covariance
.exp	* expectation
.e	- mean
.E	- covariance
.inn	* innovation
.z	- mean
.Z	- covariance
.iZ	- inverse covariance
.MD2	- squared Mahalanobis distance
.pApp	* predicted appearance
.cApp	* current appearance
.sc	* matching quality score
.vis	* lmk is visible from sensor
.measured	* lmk has been measured by the feature scanner
.matched	* lmk has been matched by the feature matcher
.updated	* lmk has been updated in SLAM
.H_r	* Jacobian of observation function wrt robot pose
.H_s	* wrt sensor
.H_l	* wrt landmark

Figure 7: The Obs structure array with its arborescence.

2.2 Simulation data

This toolbox also includes simulated scenarios. We use for them the following objects, that come with 6-letter names to differentiate from the SLAM data:

SimRob: Virtual robots for simulation.

SimSen: Virtual sensors for simulation.

SimLmk: A virtual world of landmarks for simulation.

SimObs: A virtual sensor capture, equivalent to the raw data of a sensor.

The simulation structures **SimXxx** are simplified versions of those existing in the SLAM data. Their arborescence is much smaller, and sometimes they may have absolutely different organization. It is important to understand that none of these structures is necessary if the toolbox is to be used with real data.

It follows a reproduction of the arborescences of the principal simulation data structures.

SimRob(r)	Simulated robot structure, containing:
.rob	* index in SimRob() array
.id	* robot id
.name	* robot name
.type	* robot type
.motion	* motion model
.sensors	* list of installed sensors
.frame	* frame structure, containing:
.x	- 7-vector, position and orientation $x = [t;q]$
.t	- position
.q	- orientation quaternion
.R	- rotation matrix, $R = q2R(q)$
.Rt	- transposed R
.it	- inverse position, $it = -Rt*t$
.iq	- inverse quaternion, $iq = q2qc(q)$
.Pi	- PI matrix, $Pi = q2Pi(q)$
.Pc	- conjugate PI matrix, $Pc = pi2pc(Pi)$
.vel	* velocity structure, containing
.x	- 6-vector, linear and angular velocities

Figure 8: The SimRob structure array with its arborescence.

```

SimSen(s)      Simulated Sensor structure, containing:
  .sen          * index in SimSen() array
  .id           * sensor id
  .name         * sensor name
  .type        * sensor type
  .robot        * robot it is installed to
  .frame        * frame structure, containing:
    .x          - 7-vector, position and orientation  $x = [t;q]$ 
    .t          - position
    .q          - orientation quaternion
    .R          - rotation matrix,  $R = q2R(q)$ 
    .Rt         - transposed R
    .it         - inverse position,  $it = -Rt*t$ 
    .iq         - inverse quaternion,  $iq = q2qc(q)$ 
    .Pi         - PI matrix,  $Pi = q2Pi(q)$ 
    .Pc         - conjugate PI matrix,  $Pc = pi2pc(Pi)$ 
  .par          * sensor parameters
    .k          - intrinsic params
    .d          - distortion vector
    .c          - correction vector
    .imSize     - image size

```

Figure 9: The SimSen structure array with its arborescence.

SimLmk	Simulated landmarks structure, containing:
.ids	* N-vector of landmark identifiers
.points	* 3-by-N array of 3D points
.lims	* limits of playground in X, Y and Z axes
.xMin	- minimum X coordinate
.xMax	- maximum X coordinate
.yMin	- minimum Y coordinate
.yMax	- maximum Y coordinate
.zMin	- minimum Z coordinate
.zMax	- maximum Z coordinate
.dims	* dimensions of playground
.l	- length in X
.w	- width in Y
.h	- height in Z
.center	* central point
.xMean	- central X
.yMean	- central Y
.zMean	- central Z

Figure 10: The SimLmk structure with its arborescence.

SimObs(s)	Simulated observation structure, containing:
.sen	* index to sensor in Sen()
.ids	* n-vector of measured ids
.points	* m-by-n array of measured points

Figure 11: The SimObs structure array with its arborescence. m is the dimension of the measurement space. For vision, we have $m = 2$.

2.3 Graphics data

This toolbox also includes graphics output. We use for them the following objects, which come also with 6-letter names:

MapFig: A 3D figure showing the world, the robots, the sensors, and the current state of the map.

SenFig: One figure for each sensor, visualizing its *measurement space*.

The graphics structures **XxxFig** contain data related to the graphical output. There is one figure for the 3D map, and one figure for each one of the sensors, showing the measurement space (Figs. 12, 13, 14).




Figure 12: The set of figures. The structures **MapFig** and **SenFig(s)** contain the handles to all graphics objects drawn.

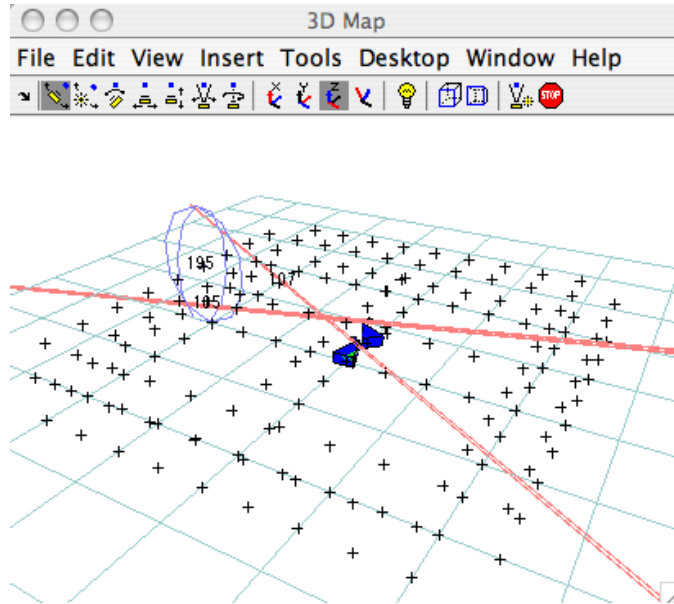


Figure 13: The 3D map figure.

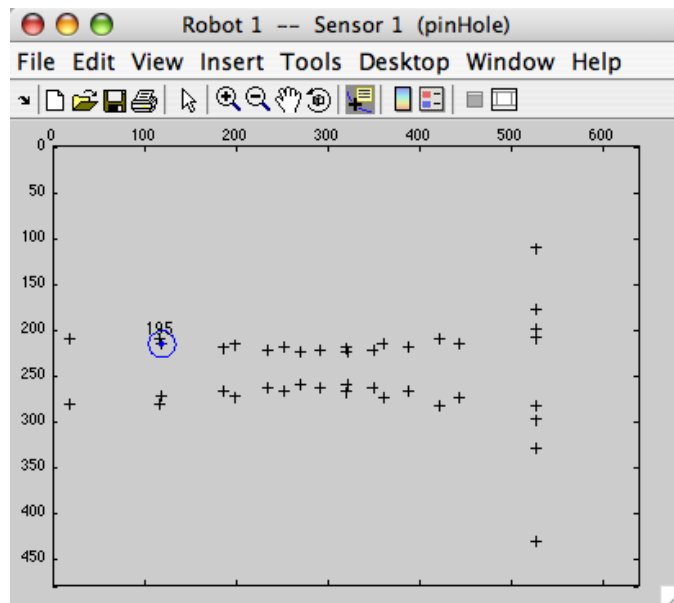


Figure 14: A pin-hole sensor view figure.

It follows a reproduction of the arborescences of the principal graphics structures.

```
MapFig          Map figure structure, containing:
  .fig          * figure number and handle
  .axes         * axes handle
  .ground       * handle to floor object
  .simRob       * array of handles to simulated robots
  .simSen       * array of handles to simulated sensors
  .simLmk       * handle to simulated landmarks
  .estRob       * array of handles to SLAM robots
  .estSen       * array of handles to SLAM sensors
  .estLmk       * handles to SLAM landmarks, containing:
    .mean       - array of handles to landmarks means
    .ellipse    - array of handles to landmarks ellipses
    .label      - array of handles to landmarks labels
```

Figure 15: The MapFig structure with its arborescence.

```
SenFig(s)       Sensor figure structure, containing:
  .fig          * figure number and handle
  .axes         * axes handle
  .raw          * handle to raw data
  .measure      * array of handles to landmarks measurements
  .ellipse      * array of handles to landmarks ellipses
  .label        * array of handles to landmarks labels
```

Figure 16: The SenFig structure array with its arborescence.

2.4 Plain data

The structured data we have seen so far is composed of chunks of lower complexity structures and plain data. This plain data is the data that the low-level functions take as inputs.

For plain data we mean:

logicals and scalars: Any Matlab scalar value such as `a = 5` or `b = true`.

vectors and matrices: Any Matlab array such as `v = [1;2]`, `w = [1 2]`
`c = [true false]` or `M = [1 2;3 4]`.

character strings: Any Matlab alphanumeric string such as `type = 'pinHole'`
or `dir = '%HOME/temp/'`.

frames: Frames are Matlab structures that we created to store data belonging to 3D frames (see Fig. 17 for an instance of the `frame` structure; type `help frame` at the Matlab prompt). We do this to avoid having to compute multiple times rotation matrices and other frame-related constructions.

A frame is specified by a 7-vector `frame.x` containing translation vector and an orientation quaternion (type `help quaternion` at the Matlab prompt). This is the essential frame information. All other fields are created or updated using the `updateFrame` function.

<code>frame</code>	Frame structure, containing:
<code>.x</code>	* the state 7-vector
<code>.t</code>	* translation vector, <code>t = x(1:3)</code>
<code>.q</code>	* orientation quaternion, <code>q = x(4:7)</code>
<code>.R</code>	* rotation matrix, <code>R = q2R(q)</code>
<code>.Rt</code>	* transposed R, <code>Rt = R'</code>
<code>.it</code>	* inverse position, <code>it = -Rt*t</code>
<code>.iq</code>	* inverse or conjugate quaternion, <code>iq = q2qc(q)</code>
<code>.Pi</code>	* PI matrix, <code>Pi = q2Pi(q)</code>
<code>.Pc</code>	* conjugate PI matrix, <code>Pc = q2Pi(iq)</code>

Figure 17: The `frame` structure and its fields.

3 Functions

The SLAM toolbox is composed of functions of different importance, defining three levels of abstraction. They are stored in subdirectories according to their field of utility. There are two particular directories: `HighLevel`, with two scripts and a limited set of high-level functions; and `InterfaceLevel`, with a number of functions interfacing the high level data with the low-level library. All other directories contain low-level functions.

3.1 High level

The high level scripts and functions are located in the directory `%SLAMtoolbox/HighLevel/`.

There are two main scripts that constitute the highest level:

`universalSLAM.m`: the main script. It initializes all data structures and figures, and performs the temporal loop by first simulating motions and measurements, second estimating the map and localization (the SLAM algorithm itself), and third visualizing all the data.

`userData.m`: a script containing the data the user must enter to configure the simulation. It is called by `universalSLAM.m` at the very first lines of code.

High-level functions exist to help initializing all the structured data. They are called by `universalSLAM` just after `userData`:

```
createSLAMstructures()  
createSimStructures()  
createGraphicsStructures()
```

Finally, other high-level functions exist for creating all graphics figures. They are called by `createGraphicsStructures.m`:

```
createMapFig()  
createSenFig()
```

3.2 Interface level

The interface level functions are located in the directory `%SLAMtoolbox/InterfaceLevel/`.

The interface level functions interface the high-level scripts and structured data with the low-level functions and the plain data. These functions serve three purposes:

1. Check the type of structured data and select the appropriate methods to manipulate them.
2. Split the structured data into smaller parts of plain data.
3. Call the low-level functions with the plain data (see Section [2.4](#)), and assign the outputs to the appropriate fields of structured data.

A good example of interface function is `SimMotion.m`, whose code is reproduced in Fig. [18](#).

```

function Rob = SimMotion(Rob, Con, dt)

% SIMMOTION Simulated robot motion.
% Rob = SIMMOTION(Rob, Con, DT) performs one motion step
% to robot Rob with control signals Con, following the
% motion model in Rob.motion. The time increment DT is
% used only if the motion model requires it, but it has
% to be provided because SIMMOTION is a generic method.
%
% See also CONSTVEL, ODO3, UPDATEFRAME.

switch Rob.motion      % check robot's motion model

    case 'constVel'

        Rob.state.x = constVel(Rob.state.x,Con.u,dt);
        Rob.frame.x = Rob.state.x(1:7);
        Rob.vel.x    = Rob.state.x(8:13);
        Rob.frame    = updateFrame(Rob.frame);

    case 'odometry'

        Rob.frame    = odo3(Rob.frame,Con.u);

    otherwise
        error('??? Unknown motion model ''%s''.',Rob.motion);
end

```

Figure 18: The interface function `SimMotion.m`. Observe that (1) the interface function checks data types and selects different low-level functions accordingly; (2) the structures are split into chunks of plain data before entering the low-level functions; (3) only `frame.x` is modified by the low-level motion functions in `case 'constVel'`, and `updateFrame()` is called afterwards; (4) the odometry function already performs frame update; (5) there is an error message for unknown motion models.

3.3 Low level library

There are different directories storing a lot of low-level functions. Although this directory arborescence is meant to be complete, you are free to add new functions and directories (do not forget to add these new directories to the Matlab path). The only reason for these directories to exist is to have the functions organized depending on their utility.

The toolbox is delivered with the following directories:

DataManagement/	Certain data manipulations
DetectionMatching/	Features detection and matching
EKF/	Extended Kalman Filter
FrameTransforms/	Frame transformations
rotations/	Rotations (inside FrameTransforms/)
Graphics/	Graphics creation and redrawing
Kinematics/	Motion models
Math/	Some math functions
Observations/	Observation models
Simulation/	Methods exclusive to simulation
Slam/	Low-level functions for EKF-SLAM

The functions contained in this directories take plain data as input, and deliver plain data as output.

4 Developing new models